

INF-2203

02- Starting OS and basic I/O


John Markus Bjørndalen
2026

Bootstrapping on an old PC

- Computer turns on in **16-bit mode** and starts running BIOS
- BIOS loads a "Boot Loader" from a designated storage device (typically your first hard disk)
 - The boot loader is a tiny piece of code that has one task: find and load the operating system kernel
 - After the kernel is loaded, the boot loader jumps to a predefined start location in the kernel and the operating system is running
 - The operating system can choose to use the BIOS to handle I/O, but can also choose to interface directly with hardware
- This is similar to CP/M, but adds the boot loader
- To get to 64-bit mode, the operating system has to (see next foil)

Historical note: CP/M and BIOS

- Basic Input/Output System (BIOS)
 - The portability layer of CP/M
 - Porting to new computers: in principle, only the BIOS would need to be specific for the computer.
- Complexity
 - The BIOS must be mapped into addr 0 from the start (CPU starts executing there)
 - The operating system stores information from addr 0
 - => need to move the BIOS before executing the operating system
 - One option: let the BIOS
 - boot computer
 - load CP/M.
 - Then the relevant part of the BIOS can be copied to higher addresses before switching out the BIOS ROM.



Bootstrapping a modern PC - implications

Implications for our OS kernel

- The old inf-2201 OS no longer boots on modern hardware
- UEFI takes care of initialisation: our kernel can start directly from 64-bit mode
 - Security mechanisms may cause issues
- Newer hardware is more complicated and may require more drivers (ACPI, ...)
- May need to consider switch to simpler architectures

Bootstrapping a newer PC through legacy mode

- Here be dragons:
- Load and start 16-bit operating system using BIOS
 - Real mode
 - Limited address space
 - Segmented memory: 16 bit addr == 64KB segments. Combine with segment register to create 20-bit addr (1MB).
 - Kernel switches to 32-bit mode. Rough description:
 - Protected mode (introduced with 80286)
 - Historical baggage: need to fiddle with keyboard driver to enable pin 20 on the address bus
 - Set up memory mapping (more about this later in the course): global descriptor table (GDT).
 - Enable 32-bit by setting Protection Enable bit in control register 0 (CR0).
 - Execute "long jump" (ljmp)
 - Kernel can now switch to 64-bit mode:
 - Long mode
 - https://wiki.osdev.org/Setting_Up_Long_Mode
 - Problems:
 - Complicated due to old baggage (backwards compatibility) – it's actually much more convoluted than the description above
 - Modern computers are dropping support for 16 and 32-bit mode:
 - Booting directly using UEFI instead of BIOS
 - Future: dropping support for 16 and 32-bit mode:
 - <https://www.intel.com/content/www/us/en/developer/articles/technical/embracing-future-simplified-architecture.html>
 - Or maybe not.. see updated note above on X86s

Bootstrapping a modern PC – the modern way

Slightly simplified

- UEFI Firmware (ROM/something) initially starts booting a computer
- The firmware looks for an EFI boot/system partition and loads a boot loader
 - The bootloader initializes computer and prepares to load an operating system
 - Bootloaders can provide menus or escape into modes for configuring the computer
- In this course, we'll use the GRUB bootloader
 - Going into details will take too much time of the course
 - Provides a boot menu
 - Locates and loads the operating system kernel and
 - See multiboot protocol for how Grub
 - Inspects the kernel binary for requests
 - Prepares the computer to correspond to requests (ex: select graphics mode)
 - Provides information to the kernel before handing over to the kernel
- If this gets complicated, don't worry. We will ignore most of it when developing the OS!

Kernel as a "program launcher"

Topics

- Basic I/O
 - Memory mapped – start with how memory reads/writes are I/O and how it's done
 - Separate I/O instructions (similar, but tagged I/O and separate addr space effectively)
- Multiboot protocol
 - Why do we need something like this?
 - Specifying what the kernel wants
 - Scanning the info the kernel gets
- Starting and running a program from the kernel
 - Need to find the files
 - how to read cpio "fs"
 - VFS

I/O – let's start with old computers

Example computer: Mycro-1

- Norwegian computer from 1974
 - Similar age to the Altair
- The first were Intel 8080 based (DIM-1001)
 - Said to be the first single-board Microcomputer (it had cpu, rom, memory and some I/O on a single card)
- Later an upgraded Z80 CPU card came (DIM-1003) that we will look at
-
- Some info:
<https://en.wikipedia.org/wiki/MYCRO-1>



Opening the computer



DIM-1001 – CPU board (Intel 8080)



DIM-1003 CPU Card (Z80)



DIM-1016 64 KB of dynamic RAM

- 16-bit addressing
 - => 64KB max



Accessing memory

- Simplified "Z80-inspired" memory read (real version is slightly more complicated)
 - CPU sets address pins on CPU and asserts /MREQ (memory request pin set low) and /RD (read pin)
 - Any memory device decodes address pins to determine if it is responsible
 - Responsible chip(s) looks up memory cells and sets content on shared data pins
 - Responsible chip(s) asserts "data ready" and CPU can copy in content from data pins
- Writes similar, just flip the /RD bit and let the CPU set the data pins while responsible chips reads the data pins



I/O - v1 – memory mapped I/O

- Observation: RAM or ROM are just chips that talk to CPU over a bus using signaling.
- Same type of reads/writes can be used by other devices, but interpreted in a different way
 - Serial port device could interpret writes to addr X and Y as commands to the serial port device
 - Writes to addr X could be writes to control register
 - Writes to addr Y could be data bytes you want to send over the serial port
- The CPU doesn't need to know what the I/O device is, only that it can read and write bytes to and from the given addresses
- The 6502 processor use memory mapped I/O
 - Simple programming, sacrifices a little bit of memory space




I/O - v2 – I/O ports

- Intel 8080 and Z80 used the same address + data pins for I/O but added another mode of communication: I/O ports.
 - Set /MREQ high in a request, signalling that the CPU is not doing a memory request
 - Memory chips ignore the operation
 - I/O port devices (like the Z80 serial port) then responds to the read/write operations
 - You could almost think of it like an extra address pin adding an I/O address space
- Modern Intel processors (x86/x64) still have the I/O ports. We will use them in the course




Multiboot protocol

- Methods for kernel to request things from bootloader
 - Setting up options in the bootloader config (like grub config), but that's done by sysadmins and not kernel
 - Specify options in the kernel binary that the bootloader can find
- The second option requires a method/protocol so the kernel can put information
 - A place the bootloader can find it
 - In a format the bootloader can understand
- This is part of the multiboot protocol




Multiboot protocol

- Time for some code
 - grub-multiboot2/test-kern
 -



Multiboot protocol

- How does the kernel know what the bootloader gave it?
- Bootloader fills in a data structure in memory that describes things like
 - State of computer (memory maps...)
 - Where loaded files are located
 - ...
- Before starting the kernel, the address of that data structure is written to a register
- The first thing a kernel does is to store that address somewhere and then start parsing the data structure
- Back to the example code



Multiboot protocol, resources

For more information


- We use the multiboot protocol for our kernel, so you can have a look there
- <https://www.gnu.org/software/grub/manual/multiboot/multiboot.html>
- <https://wiki.osdev.org/Multiboot>

Bundling files for the kernel – initrd

- Initrd = Initial RAM disk
- A binary file that contains a simple file system
 - One of the simplest in common use is "cpio" (we will get back to this)
- Bootloader loads it into memory when the kernel is loaded, then the bootloader tells the kernel where the initrd is (using multiboot)
- Typically contains support files that the kernel needs to boot up
 - Drivers
 - Scripts for initializing and setting up a system
 - Initial binaries and programs
- In this course, we will use a cpio based initrd to contain things like programs we want to start running

The CPIO file format (simplified)

- There are multiple versions and options for headers, but we will look at a simple version
- The basic format is as follows:
 - Files are appended to the binary file with the following:
 - A header that describes one file, with information such as file name, size, modes, user ids, file types etc.
 - The binary file added just after the header
 - Similar in idea to a tar or zip file (except zip files have indexes to where files are)
- To find a file, you scan through the file until you find a matching file name.
- Code example: ramfstst
 - You will get a cleaner example in the P1 source code




Using a CPIO / initrd file as initial file system

- You can view any archive file like cpio as a file system
- To use it from programs, you need functions for
 - Locating files
 - Opening files
 - Reading files
 - We're ignoring file writes for now
- We could provide "cpio filesystem" functions to programs
 - Would require all programs to know about all file systems we want to use, and figure out which to use

Using a CPIO / initrd file as initial file system - VFS

- Alternative approach: provide an abstraction to any filesystem
 - Think OOP or abstract data types
 - Provide a file system data structure that describes the file system and operations on it (opening files, finding files etc)
 - First part of a file path could determine filesystem, then look inside fs
 - In our operating system: VFS
 - lib/drivers/vfs_fs.h and vfs_file.h
 - Mike?

(over)simplified example of VFS idea



TODO

- ELF format
- Loading and starting programs (jump or function call)